What Are We Measuring? On Digital Twinning, Behavioral Prediction, and the Epistemology of AI Psychological Assessment

Peter Golub, March 2026

Abstract

The rapid convergence of large language models (LLMs), digital phenotyping, and computational psychometrics has opened a new frontier in personality assessment: the extraction of psychological constructs from natural language without explicit self-report instruments. Recent empirical work—most notably Wright et al. (*Nature Human Behaviour*, 2026), demonstrating r = .88 convergence between zero-shot LLM scoring and self-reported Big Five traits—validates the foundational premise that personality is sedimented in everyday language and computationally recoverable. Yet this validation raises an epistemological question that the field has not adequately addressed: when an AI system produces a psychological evaluation from conversational data, what is the ontological status of the output? This paper argues that AI-generated personality assessments are neither measurements of stable internal traits nor predictions of future behavior, but cognitive models—computational schemas of individual behavioral tendencies with specifiable accuracy bounds and structural limitations. Drawing on Park's (2025) generative agent simulations (the fidelity ceiling: 85% attitudes, 80% personality, 66% novel behavior), Barrett's (2017) theory of constructed emotion, Heuer's (1999) analysis of cognitive bias in intelligence assessment, Ariely's (2012) research on universal self-deception, and Paul's (2004) history of personality testing, we propose a framework for honest psychological assessment that reports its own boundary conditions, models person-environment interactions rather than isolated traits, and implements structured self-correction through competing hypotheses. We introduce the concept of the psychological digital twin—a longitudinal, context-sensitive, episodically grounded cognitive model that occupies a defined position on the spectrum between biological twinning (deterministic, high-fidelity) and complete psychological simulation (irreducibly complex, structurally bounded). We conclude with architectural recommendations for assessment platforms seeking to build instruments that know where they go blind.

Keywords*: personality assessment, large language models, digital twin, psychometrics, cognitive modeling, fidelity ceiling, episodic memory, Big Five*

1. Introduction: The Cult and the Calculus

The history of personality assessment is a history of overconfidence. Annie Murphy Paul's The Cult of Personality Testing (2004) documents a century of instruments—phrenology, Rorschach, MMPI, Myers-Briggs, the Big Five—each of which attacked the previous generation's reductionism and then performed its own. Phrenologists measured skulls and claimed to read character. Rorschach administered inkblots and claimed to access the unconscious. The MMPI administered 567 true-false questions and claimed to diagnose psychopathology. Myers-Briggs sorted four billion dollars' worth of humanity into sixteen types. The Big Five reduced the space to five orthogonal dimensions and called it progress. Each instrument survived not primarily through scientific validation but through institutional embedding—adoption by employers, militaries, courts, and therapists who needed a number, any number, to justify a decision about a person (Paul, 2004).

Meanwhile, the science of psychology itself remains in the throes of the replication crisis, with widespread failures to reproduce published results calling into question the empirical foundations on which these instruments rest. And yet we all know that personalities exist. Different people are less or more likely to act in specific ways, and our relationships are defined by how we predict—often accurately—what people will do and say. We build cognitive models of people every day. The question is whether a machine can do it better.

The answer, as of 2026, is: partly. Recent advances in large language models have demonstrated that commercially available AI systems can extract personality-relevant information from natural language with a fidelity that would have seemed impossible five years ago. Wright et al. (2026) showed that seven commercially available LLMs, prompted to score Big Five traits from brief open-ended narratives with no model-specific training, achieved convergence with self-reports comparable to or exceeding meta-analytic benchmarks for family and friend ratings. The lexical hypothesis—the foundational premise that the most important attributes of personality are sedimented in everyday language—is no longer speculative. It is empirically demonstrated at the highest level of peer-reviewed science.

But demonstration of capability is not the same as understanding of what the capability produces. When an LLM reads a conversation and generates a personality profile, what has it done? Has it measured the person, the way a thermometer measures temperature—reading a property that exists independently of the instrument? Has it predicted behavior, the way a weather model forecasts rain—probabilistically, with specifiable error? Or has it constructed a representation, the way a portrait painter renders a face—capturing something real but shaped irreversibly by the method of capture?

This paper argues for the third interpretation, and proposes that the appropriate conceptual framework is not measurement or prediction but modeling—specifically, the construction of a cognitive digital twin: a computational schema of an individual's behavioral tendencies, linguistic patterns, emotional registers, and attitudinal positions, constructed from their conversational data, capable of generating predictions within specifiable domains, and explicitly degraded outside them. We develop this framework through a synthesis of recent empirical findings in zero-shot personality assessment, computational simulation of human behavior, the psychology of self-deception, and the epistemology of intelligence analysis. We show that the framework resolves persistent tensions in the field—between trait and situation, between accuracy and honesty, between commercial viability and scientific rigor—and yields specific architectural recommendations for platforms seeking to build instruments that know where they go blind.

2. The State of the Science: What LLMs Can and Cannot Do

The empirical case for LLM-based personality assessment rests on converging evidence from multiple independent research programs. Wright et al. (2026) administered personality inventories to seven commercially available LLMs (including GPT-4, Claude, Gemini, Grok, and LLaMA), prompting each to score individual items from the NEO-FFI and IPIP-NEO inventories as if it were the participant who produced the narrative. Using two distinct samples—spontaneous streams of thought ($N = 60$) and daily video diaries ($N = 108$)—the authors demonstrated that the average LLM-generated trait score achieved convergence with self-reports comparable to meta-analytic family and friend ratings. The effect held across both samples despite radically different data collection contexts, and the LLM-derived scores successfully predicted daily behaviors, mental health outcomes, and psychiatric diagnoses.

Three findings from Wright et al. deserve particular emphasis. First, no single LLM consistently outperformed the others; the average across models was more reliable than any individual model, paralleling the well-established wisdom-of-crowds effect in human judgment. Second, the LLMs dramatically outperformed traditional natural language processing methods (LIWC) and matched or exceeded bespoke machine learning models (MyPersonality) that were trained on millions of social media posts. Third, the convergence was achieved from remarkably brief samples—approximately 15 minutes of spontaneous thought or 5 to 7 daily diary entries were sufficient to approximate the full-sample effect sizes. These findings suggest that personality information is not merely present in language but is densely encoded and recoverable from small samples by systems with sufficient contextual understanding.

Serapio-Garcia et al. (2025) established a complementary result: that LLMs not only can assess personality but can be made to exhibit reliable personality profiles themselves. Administering standard inventories to 18 LLMs, the authors found that larger, instruction-fine-tuned models exhibited psychometric reliability and validity comparable to human respondents, and that their synthetic personalities could be deliberately shaped through structured prompting. This bidirectional finding—LLMs can both read and simulate personality—has implications for assessment that extend beyond scoring.

Hussain et al. (2025) addressed a practical constraint: the input length limitation of standard transformers. By combining a fine-tuned RoBERTa model with a recurrent neural network for context aggregation, the authors achieved competitive personality prediction from long-form life narratives (averaging over 2,000 tokens) at a fraction of the computational cost of frontier models. The hybrid architecture outperformed larger models including Longformer and LLaMA, demonstrating that domain-specific architectural decisions can substitute for scale. For assessment platforms operating under computational constraints, this result indicates that the engineering challenge is tractable.

Taken together, these findings validate the core premise: personality is computationally recoverable from natural language. The correlation coefficients are real. The predictive validity is real. The practical applicability—cheap, fast, scalable, requiring no specialized training—is real.

What the findings do not validate is any specific claim about what the recovered constructs represent. This requires a different kind of evidence.

3. The Fidelity Ceiling

Joon Sung Park's dissertation (2025) on generative agent simulations of human behavior provides the most rigorous empirical map of the boundary between what language-model-based cognitive simulation can and cannot do. Park recruited 1,052 Americans, conducted two-hour semi-structured interviews, embedded the transcripts as episodic memory in generative agents equipped with reflection and planning modules, and tested the agents against the real humans' responses across 164 outcome areas spanning politics, personality, economics, and social behavior.

The results define a fidelity ceiling:

- 85% accuracy on attitudinal questions
- 80% correlation on personality assessments
- 66% accuracy on novel behavioral scenarios

These numbers are not benchmarks to be exceeded. They are a structural map of the relationship between language and action. The first two confirm that LLMs can extract stable

psychological constructs from conversational data with impressive fidelity—agents built from interview transcripts correlated with humans' own self-reported personality at r = .80, closely matching Wright et al.'s independent finding. But the third number reveals the ceiling's nature: when agents were placed in genuinely novel game-theoretic experiments (dictator games, trust games, public goods games), their predictions diverged from real human behavior by 19 percentage points. The agents overestimated cooperation, overestimated generosity, and produced smoother, more prosocial behavior than the actual humans, who were messier, more selfish, more context-sensitive, and more irrational.

Park's architectural contribution is equally important. His demographic agents—given only age, gender, race, and location—predicted at 76%. His persona agents—given a short biographical paragraph—predicted at 78%. His interview agents—given the full transcript with expert reflections—predicted at 85%. The jump from 78% to 85% comes not from a smarter model but from richer data structured as memory. The interview transcript functions as episodic memory; the expert reflections function as semantic memory. Together, they give the model a particular person rather than a statistical average. This is the architectural lesson: the path to higher fidelity is persistent memory, not model intelligence.

Park is also honest about a finding that many in this field prefer to ignore: his agents replicate the biases of self-reports, not just their signal. Where the General Social Survey is known to produce socially desirable responses—on questions about race, charitable giving, civic participation—the agents reproduce the same inflation. They are mirrors, not x-rays. They reflect what people say about themselves, not necessarily what people are.

4. The Fudge Factor: Why Organic Data Is Not Unfiltered Truth

The Behavioral Veracity Theory of Personality (BVT-P), as articulated in Hypatia's foundational documentation, rests on a distinction between formal self-reports—surveys, structured assessments, the artificial pragmatics of the testing situation—and organic conversational data: Discord logs, text messages, unstructured social interaction. The theory claims that formal self-reports are contaminated by social desirability bias, while organic conversation reveals the authentic person because the speaker is not performing for an assessor.

Dan Ariely's research on dishonesty (2012) destabilizes the second half of this claim. Ariely's central finding is not that people lie strategically. It is that everyone lies a little, all the time, and that this micro-dishonesty is not noise—it is structural. He calls it the fudge factor: the amount of self-enhancing distortion a person can engage in while still maintaining their self-concept as an honest person. The fudge factor does not respond to rational incentives (probability of getting caught, magnitude of punishment). It responds to narrative conditions—whether the cheating involves tokens or cash, whether moral priming has occurred, whether the person has recently engaged in creative thinking. Creative individuals, Ariely demonstrates, cheat more—not because they are less moral but because they are better at constructing justifying narratives.

The fudge factor is not a property of formal assessment contexts. It is a property of human cognition. It operates in every conversation, in every casual self-description, in every anecdote about how one handled a conflict at work. The speaker is not lying to their friend on Discord. They are telling themselves a story in which they were calm and professional, and the story arrived pre-edited by the machinery of self-regard—unconsciously, automatically, at a level below narrative awareness.

This is, in functional terms, the same mechanism Lisa Feldman Barrett describes in her theory of constructed emotion (2017). Barrett argues that emotions are not biologically hardwired reactions detected by the brain but predictions constructed from prior experience, current interoceptive signals, and cultural concepts. The psychiatrist who looks at a patient and detects sadness is not

reading the patient's interior. The psychiatrist is constructing an interpretation and experiencing it as perception. The construction happens so fast that the constructor experiences it as transparent access to the truth. Ariely's fudge factor operates through the same mechanism: people generate narratives about their internal states, and the generation is automatic, unconscious, and systematically biased in the direction of self-enhancement.

Park's 19-point gap between attitudinal accuracy (85%) and behavioral accuracy (66%) is, in Ariely's framework, the fudge factor's footprint in language. The interview transcripts were already fudged—structurally, not fraudulently. The agents learned from the fudged versions and reproduced the fudge faithfully.

The implication for assessment platforms is direct: organic conversational data is not unfiltered truth. It is self-report that the speaker does not recognize as self-report. This makes it harder to game deliberately—a genuine advantage—but it does not make it more accurate. The distinction between "hard to game" and "authentic" is the distinction Hypatia must internalize at the architectural level.

5. The Digital Twin: A Framework for Honest Assessment

If AI-generated personality evaluations are neither measurements of stable traits nor predictions of behavior, what are they? We propose that they are cognitive models—and that the appropriate conceptual framework is the digital twin.

The digital twin concept originated in aerospace engineering (NASA, early 2000s) and has expanded into biomedicine, manufacturing, and healthcare. The Digital Twin Consortium defines a digital twin as an integrated, data-driven virtual representation of a real-world entity that features synchronized interaction at a specified frequency and fidelity. Unlike static models, digital twins are bidirectional: they are continuously updated with new data from the original system, and their predictions can be tested against that system's actual outputs.

Stanford Medicine's 2025 study of the mouse visual cortex illustrates the concept at its most rigorous: an AI model trained on over 900 minutes of neural recordings can predict how thousands of individual neurons will fire in response to novel stimuli, run millions of virtual experiments in hours rather than years, and reveal structural insights (such as stimulus-preference-based connectivity) that escaped direct observation. This is a genuine digital twin. It earns the name because it satisfies three conditions: it models a system with deterministic dynamics constrained by physical law; it is bidirectionally synchronized with the original system; and its accuracy envelope is specifiable.

Building a digital twin of a person's psychology satisfies none of these conditions fully—but it partially satisfies all three, and the partial satisfaction is measurable. The person's psychology is not deterministic, but it exhibits statistical regularities (the Big Five are real dimensions of variance). The twin cannot be bidirectionally synchronized in real time, but it can be updated with new conversational data as it becomes available. The accuracy envelope is not as tight as neuronal firing prediction, but it has been empirically measured: approximately 85% on attitudes, 80% on personality, 66% on novel behavior.

We propose a spectrum of twinning fidelity with three reference points:

- At one end: the biological twin (Stanford mouse cortex). Deterministic dynamics, physical constraints, specifiable accuracy. This is the Outer Self—biology governed by physics.

- In the middle: the behavioral twin (Oracle Data Cloud). Thirty thousand data categories per individual, five billion profiles, exceptional accuracy at predicting routine consumer

behavior. Oracle knows what you will buy next because the best predictor of future behavior is past behavior. But Oracle's profile is psychologically empty. It tracks behavioral exhaust—clicks, purchases, locations—without modeling the cognitive architecture that generates the behavior. It knows what you did, not why. It has a model of the footprints, not the mind.

- At the other end: the complete Inner Self. Narrative identity, moral reasoning, emotional volatility, fluid intelligence, the capacity to surprise yourself. This cannot be twinned—not now, and not with any foreseeable architecture. The irreducible gap between model and mind is not an engineering problem. It is a philosophical one. The mind is a system that changes its own rules, sometimes in response to the very models that try to capture it.

Hypatia's cognitive twin lives in the middle of this spectrum. It does more than Oracle—it builds a cognitive model, not just a behavioral trail. It does less than the complete Inner Self—it captures what language reveals and misses what language conceals. Its accuracy has a ceiling, and the ceiling has been measured. Its predictions are probabilistic, context-dependent, and degraded outside the domains in which they were validated.

6. $B = f(P, E)$: Personality as Function, Not Trait

Kurt Lewin's equation—*Behavior = f(Person, Environment)*—is the oldest and most ignored insight in psychology. Every personality test since Woodworth's Personal Data Sheet (1917) has attempted to isolate P from E and measure P as a stable quantity. Every test has partially failed, because P is not independent of E. The same person is different at work and at home, different at twenty-five and at forty, different on the Discord server and on the Slack channel. What we call personality is not a fixed property but a pattern of regularities in the $P \times E$ interaction.

The Big Five captures genuine statistical regularities. But it strips away exactly the contextual information that determines how those regularities manifest in any particular situation. The digital twin framework points toward a different product: not a profile that says you are 0.78 neurotic, but a model that says: in professional contexts with structured task demands, this person displays low hedging and consistent emotional regulation; in intimate conversations, the same person displays significantly higher emotional volatility and hedging frequency three standard deviations above their professional baseline; the gap has been widening over the last four months.

That is not a score. That is $B = f(P, E)$ rendered as a dynamic, context-sensitive, longitudinal model. It captures the interaction between person and environment, not the person in isolation. And it is the product that no one else is building—not because the technology is unavailable, but because no one else has the epistemological framework to know what they should be building with it.

7. The Analyst's Posture: Bias, Competing Hypotheses, and Epistemic Humility

Richards Heuer's Psychology of Intelligence Analysis (1999) was written for CIA analysts, but its central insight applies with uncomfortable directness to computational personality assessment: cognitive biases are not bugs in an otherwise rational system but architectural features of human cognition. Confirmation bias, anchoring, mirror-imaging—these are not errors that training can eliminate. They are properties of pattern-matching systems, biological or computational, that operate on incomplete data.

When an assessment platform prompts a model to evaluate a conversation for indicators of neuroticism, the model will find indicators of neuroticism. Not because it fabricates evidence, but because the prompt has set the prior. The prompt is the hypothesis. The text is the evidence. The profile is the judgment. And the judgment is shaped by the hypothesis at least as much as by the evidence.

Heuer's solution was not to eliminate bias but to structure the analytical process so that biases are visible and checkable. His Analysis of Competing Hypotheses (ACH) forces analysts to generate multiple explanations, evaluate evidence against each, and reason by elimination rather than confirmation. For psychological assessment, this means the system should not ask which personality profile best fits this text? but which personality profiles are consistent with this evidence, and which can the evidence eliminate?

Concretely, the twin should output not a single profile but a matrix of competing hypotheses: this person is high in neuroticism; this person is experiencing situational stress with a stable baseline; this person's patterns reflect cultural norms rather than individual traits; this person is performing a social identity rather than expressing an internal state. For each hypothesis, the system identifies supporting and contradicting evidence. The surviving hypothesis is the tentative conclusion—tentative, accompanied by the evidence matrix, and revisable as new data arrives.

The cognitive analogue from Heuer's most dangerous bias—mirror-imaging—also applies. When the model evaluates a speaker from a context underrepresented in its training data, it applies learned correlations that may not hold and generates confident predictions that may be systematically wrong. The model cannot detect this. Hypatia should flag every assessment with the demographic and cultural distance between the assessed person and the populations on which the psychometric frameworks were validated.

8. The Hermeneutical Lacuna and the Case for Multimodal Architecture

Purely text-based assessment creates what we term a hermeneutical lacuna—a gap in interpretive understanding where the system lacks the environmental, physiological, and non-verbal cues that constitute the majority of diagnostic information in clinical practice. A text parser misses the trembling voice, the prolonged hesitation, the physiological spikes of autonomic anxiety. Research in clinical psychology consistently demonstrates that conditions like anxiety and depression are identified substantially through vocal prosody, speech hesitation, and facial micro-expressions that text-only systems cannot access.

Furthermore, the pervasive use of AI interfaces introduces cognitive offloading: individuals habitually lean on digital platforms to recall facts and solve problems, eroding their own memory schemas. The Source Monitoring Framework (SMF) predicts that when individuals co-create text with an AI, their ability to distinguish whether an idea originated from their own cognition or was suggested by the system degrades significantly. An assessment that relies on text produced in collaboration with AI may inadvertently measure the user's regurgitation of the AI's own synthesized prompts.

The architectural response is multimodal integration. Platforms that combine text analysis with vocal feature extraction (pitch, cadence, hesitation patterns via models like Wav2Vec or HuBERT) and facial micro-expression analysis can cross-reference semantic claims against physiological reality. If a user's text suggests confidence but their vocal cadence reveals stress, the discrepancy itself is diagnostic information. This multimodal fusion directly addresses the hermeneutical lacuna—not by eliminating it, but by narrowing the gap between what the system observes and what the person experiences.

9. What the Twin Cannot Do: Honest assessment requires explicit statement of structural limitations. Three are irreducible

First, the saying-doing gap. The model is built from language. Language encodes what people are willing to say about themselves. Behavior is what happens when the social scaffolding falls away. The model predicts what the person will say with approximately 85% accuracy and what they will do with approximately 66% accuracy. The 19-point gap is not closable by any text-based architecture, because the gap is a property of the relationship between language and action, not a deficit of information.

Second, training bias invisibility. The model's personality constructs were developed on WEIRD populations and validated within specific cultural and institutional contexts. When the model evaluates a speaker from an underrepresented context, it mirror-images—generating confident predictions from correlations that may not hold. The model cannot detect its own mirror-imaging.

Third, the Galaxy problem. Park documented that human judges could distinguish AI-generated content from real content only 59% of the time. In an environment saturated with AI-generated text, the organic conversational data that assessment platforms analyze may not always come from organic humans. A personality assessment performed on a synthetic persona produces a profile of a fictional character presented as a diagnosis of a real person. This is not a hypothetical concern. Discord servers, customer service channels, and social media platforms already contain AI participants indistinguishable from humans.

10. Architectural Recommendations

The cognitive twin framework implies a product that does not yet exist in the personality assessment market: a longitudinal cognitive modeling platform with the following architectural commitments.

Persistent episodic memory. Every conversation analyzed becomes part of a growing, time-indexed archive. New assessments draw on the full history. The system tracks change over time, not just state at a moment. Park showed that the jump from 78% to 85% comes from richer memory, not smarter models.

Context-sensitive modeling. The system outputs $B = f(P, E)$: how this person behaves across contexts, how the contexts interact with the person, and how the interaction has changed over time.

Competing hypotheses. Every assessment includes an ACH-style matrix. Multiple interpretations are generated, evidence is evaluated against each, and the output includes not just the leading hypothesis but the alternatives and the evidence that would distinguish them.

Asymmetric confidence intervals. The fudge factor biases self-description toward self-enhancement. The system should report estimates with tighter bounds on the low end and wider bounds on the high end.

Explicit boundary reporting. Every assessment states: this model predicts self-reported attitudes at approximately 85%, personality-relevant linguistic behavior at approximately 80%, and behavior in novel high-stakes contexts at approximately 66%. These are the boundary conditions of the instrument.

Longitudinal fudge-factor detection. Over time, the gap between a person's self-descriptions and their described outcomes becomes observable. This gap is the fudge factor made visible across dozens of conversations.

Synthetic persona detection (Galaxy problem infrastructure). The system must verify that the conversational data was produced by a human. In an AI-saturated environment, this is a core engineering challenge that, if solved, constitutes a genuine competitive advantage.

Multimodal integration pathway. Text-only assessment should be the foundation, not the ceiling. The architecture should be designed for progressive integration of vocal, physiological, and visual channels as they become technically and ethically feasible.

11. Conclusion: Build the Twin, Report the Gap

The question "What are we measuring?" has an answer, but it is not "personality" per se. We are building a model. A computational cognitive twin that runs in parallel with the person, makes predictions about their behavior, and is continuously updated as new data arrives. The model is not the person. The model is a lossy compression of the person's linguistic output, shaped by the instrument's architecture, the biases of its training data, the pragmatics of the conversational context, and the universal human tendency to present a slightly edited version of the self to the world. The compression captures real regularities—personality is sedimented in language, and the language model can read it—but it inherits real distortions, and the distortions are systematic, documentable, and approximately quantified.

The honest version of this product—the version that reports its own boundary conditions, tracks its own biases, models the person-environment interaction rather than the person in isolation, and tells the user where the predictions are reliable and where they are not—is more valuable than the dishonest version. Not because the market currently rewards honesty. The market currently rewards clean scores, simple labels, and confident-sounding output that requires no interpretation. But the market for clean scores is the market that sustained phrenology, the Rorschach, and the MBTI—the market that keeps buying the sieve and calling it a thermometer.

Hypatia can build a different instrument. Not a better sieve. Not a smarter survey. A cognitive twin—bounded, honest, longitudinal, and more useful than any personality score precisely because it knows what it cannot do.

Build the twin. Report the gap. That is the product.

References

Ariely, D. (2012). *The (Honest) Truth About Dishonesty: How We Lie to Everyone—Especially Ourselves.* Harper.

Alsini, R., Naz, A., Khan, H. U., Bukhari, A., Daud, A., & Ramzan, M. (2024). "Using deep learning and word embeddings for predicting human agreeableness behavior." *Scientific Reports*, 14, 29875.

Barrett, L.F. (2017). *How Emotions Are Made: The Secret Life of the Brain*. Houghton Mifflin Harcourt.

Burton, J.W., et al. (2024). "How large language models can reshape collective intelligence." *Nature Human Behaviour*, 8, 1643–1655.

Heuer, R.J., Jr. (1999). *Psychology of Intelligence Analysis*. Center for the Study of Intelligence, Central Intelligence Agency.

Hussain, R., Ma, J., Khandelwal, R., Oltmanns, J., & Gupta, M. (2025). "Personality prediction from life stories using language models." arXiv preprint.

Lewin, K. (1936). *Principles of Topological Psychology*. McGraw-Hill.

Naz, A., Khan, H.U., Bukhari, A., Alshemaimri, B., Daud, A., and Ramzan, M. (2025). "Machine and deep learning for personality traits detection: A comprehensive survey and open research challenges." *Artificial Intelligence Review*, 58, 239.

Park, J.S. (2025). *Generative Agent Simulations of Human Behavior*. Doctoral dissertation, Stanford University.

Park, G., Schwartz, H.A., Eichstaedt, J.C., Kern, M.L., Kosinski, M., Stillwell, D.J., Ungar, L.H., and Seligman, M.E.P. (2015). "Automatic personality assessment through social media language." *Journal of Personality and Social Psychology*, 108(6), 934-952.

Paul, A.M. (2004). *The Cult of Personality Testing: How Personality Tests Are Leading Us to Miseducate Our Children, Mismanage Our Companies, and Misunderstand Ourselves*. Free Press.

Serapio-Garcia, G., Safdari, M., Crepy, C., Sun, L., Fitz, S., Romero, P., Abdulhai, M., Faust, A., and Mataric, M. (2025). "A psychometric framework for evaluating and shaping personality traits in large language models." *Nature Machine Intelligence*, 7, 1954-1968.

Stanford Medicine. (2025, April). "AI models of the brain could serve as "digital twins" in research." *Stanford Medicine News*.

Steyvers, M., Tejeda, H., Kumar, A., Belem, C., Karny, S., Hu, X., Mayer, L.W., and Smyth, P. (2025). "What large language models know and what people think they know." *Nature Machine Intelligence*, 7, 221-231.

Wright, A.G.C., Ringwald, W.R., Vize, C.E., Eichstaedt, J. C., Angstadt, M., Taxali, A., & Sripada, C. (2026). "Assessing personality using zero-shot generative AI scoring of brief open-ended text." *Nature Human Behaviour*. https://doi.org/10.1038/s41562-025-02389-x

Yampolskiy, R.V. (2015). *Artificial Superintelligence: A Futuristic Approach*. Chapman & Hall/CRC.